

Approximate learning of k -step invariant for probabilistic program

Arnab

September 24, 2024

Chapter 1

Approximate learning of k -step invariant

1.1

Definition 1.1.1. k -step invariant of a probabilistic program P : Given a probabilistic program P and a set S of initial states, the k -step invariant or the k -iterative closure of P with respect to S is given by the set $cl_k(S) = \{\sigma \in \{0,1\}^{|V|} \mid \sigma \text{ is reachable from } S \text{ in at most } k \text{ iterations of } P\}$. Note that $cl_k(S) = \cup_{i=0}^k \psi^i(S) = \{\sigma \in \{0,1\}^{|V|} \mid \sigma \text{ is reachable in at most } k \text{ steps from } S\}$, where ψ denotes the single-iteration input-output function for the program P .

Definition 1.1.2. Distance between a candidate and the k -step invariant: Given a probabilistic program P and a set S of initial states, the distance between a given candidate T and the k -step invariant $cl_k(S)$ is given by $d(T, cl_k(S)) = \sum_{y \in (cl_k(S) \setminus T)} \Pr_{x \sim_U S}[y \text{ is reachable from } x \text{ in at most } k \text{ iterations}] = \Pr_{x \sim_U S}[(cl_k(S) \setminus T) \text{ is reachable from } x \text{ in at most } k \text{ iterations}]$.

Assumption 1.1.1. Succinct description of the true k -step invariant: We have assumed that the true k -step invariant for a given probabilistic program P with respect to a set S of initial states can be expressed in the form of a CNF of bounded size (polynomial) in the number of variables.

Assumption 1.1.2. Sampling access to internal random variables of P : We have assumed that the probabilistic program can be allowed to run in a deterministic manner by fixing a random seed, given sampling access to the internal random variables R .

Note. Monotonicity of candidates with respect to violating transitions: An important thing to note here is that for the ease of analysis, we are only interested in the family Γ of candidates such that given any candidate $T \in \Gamma$, any k -length transition starting from S does not return back to T once it goes out of it.

Problem 1.1.3. (Additive approximation of the distance of a candidate T from k -step invariant $cl_k(S)$): Given a probabilistic program P , a set S of initial states, the number of iterations k and a candidate T , parameters $\epsilon, \delta \in (0, 1]$ output an ϵ -additive approximation of the distance $d(T, cl_k(S))$ with probability at least $1 - \delta$.

Theorem 1.1.4 (Correctness of DistEstimate). Given a probabilistic program P , a set S of initial states, the number of iterations k , a candidate T expressed as a CNF, parameters $\epsilon, \delta \in (0, 1]$, **DistEstimate** outputs an ϵ -additive estimate of the $d(T, cl_k(S))$ with probability at least $1 - \delta$. Also, **DistEstimate** requires at most $\lceil \frac{1}{2\epsilon^2} \log(\frac{2}{\delta}) \rceil$ samples from $\text{Unif}(S \times \mathcal{P}(R))$.

Algorithm 1 IsNotWitness(T, w)

- 1: Initialize $\tau \leftarrow 0$.
 - 2: $\tau \leftarrow T(w)$
 - 3: Output $\neg\tau$.
-

Algorithm 2 DistEstimate($P(V, R), S, T, \epsilon, \delta, k$)

- 1: Initialize $m \leftarrow \lceil \frac{1}{2\epsilon^2} \log(\frac{2}{\delta}) \rceil$, $S_U \leftarrow \emptyset$, $\hat{d}_{S_U} \leftarrow 0$, $\tau \leftarrow 0$.
 - 2: $S_U \leftarrow m$ iid samples from $\text{Unif}(S \times \mathcal{P}(R))$.
 - 3: **for** $i \in [m]$ **do**
 - 4: With $(x_i, R_i) \in S_U$ as initial state, run the program P for k iterations to obtain an output state y_i .
 - 5: $\tau \leftarrow \text{IsNotWitness}(T, y_i)$
 - 6: $\hat{d}_{S_U} \leftarrow \hat{d}_{S_U} + \frac{\tau}{m}$
 - 7: **end for**
 - 8: Output \hat{d}_{S_U} .
-

For a given candidate T , we can write the distance of T from the k -step invariant $cl_k(S)$ as follows:
 $d(T, cl_k(S)) = \Pr_{x \sim_U S}[(cl_k(S) \setminus T) \text{ is reachable from } x \text{ in at most } k \text{ iterations}]$.

Proof of correctness of DistEstimate:

Claim: Given a probabilistic program P , a set S of initial states, the number of iterations k , a candidate T expressed in CNF, parameters $\epsilon, \delta \in (0, 1]$, **DistEstimate** outputs an estimate \hat{d}_{S_U} of the distance $d(T, cl_k(S))$ with the following guarantees:

$$\Pr[|\hat{d}_{S_U} - d(T, cl_k(S))| \leq \epsilon] \geq (1 - \delta)$$

Moreover, the objective of **DistEstimate** is to generate a set of positive counterexamples, i.e., those states which have not been learnt by the candidate but they do actually belong to the true k -step invariant $cl_k(S)$.

Proof. Description of IsNotWitness: **IsNotWitness** takes in a CNF formula T and a given assignment w of the variables $V \in \text{supp}(T)$ and returns 1 if w is not a witness for T , and 0 otherwise.

Description of DistEstimate: Line 2 samples m states from $(S \times \mathcal{P}(R))$ uniformly at random so as to obtain a sample set $S_U = \{(x_1, R_1), (x_2, R_2), \dots, (x_m, R_m)\}$, where for each $i \in [m]$, x_i is the initial state and R_i is the fixed initial seed for the internal random bits of P .

Starting from each state defined by $(x_i, R_i); i \in [m]$, the program P is executed for exactly k iterations in Line 4. According to assumption , if for a given initial state (x_i, R_i) , the k -length transition to output state y_i goes out of the set represented by T at some j -th iteration ($j \in [k]$), y_i is guaranteed not to be a witness for the candidate T . Hence, it is sufficient to just check the output state y_i reached after k iterations.

Now, let's define the following event for each state $(x_i, R_i) \in S_U; i \in [m]$:

E_i : With (x_i, R_i) as initial state, the output state y_i reached after k iterations is not a witness for T .

Note that the collection of events $\{E_i\}_{i=1}^m$ are mutually independent. Next, we define indicator random variables for these events as $\mathbf{1}_{E_i}$. Define the statistic $\hat{d}_{S_U} = \frac{1}{m} \sum_{i=1}^m \mathbf{1}_{E_i}$. Line 6 updates the estimate \hat{d}_{S_U} based on the check for non-witness of output state y_i for the candidate T in line 5 via the subroutine **IsNotWitness**.

We observe that $\mathbb{E}[\hat{d}_{S_U}] = \Pr_{x \sim_U S}[(cl_k(S) \setminus T) \text{ is reachable from } x] = d(T, cl_k(S))$ and thus, \hat{d}_{S_U} is an unbiased estimator of the quantity $d(T, cl_k(S))$. Applying additive Chernoff bound given an error parameter $\epsilon \in (0, 1]$, we get

$$\Pr[|\hat{d}_{S_U} - \mathbb{E}[\hat{d}_{S_U}]| \geq \epsilon] = \Pr[|\hat{d}_{S_U} - d(T, cl_k(S))| \geq \epsilon] \leq 2e^{-2m\epsilon^2}$$

We want to make this probability go below a certain threshold, given by δ . Thus,

$$2e^{-2m\epsilon^2} \leq \delta \implies m \geq \frac{1}{2\epsilon^2} \log\left(\frac{2}{\delta}\right).$$

This gives us a sample complexity of $O(\frac{1}{\epsilon^2} \log(\frac{2}{\delta}))$. Thus, we can conclude that if we take atleast $\lceil \frac{1}{2\epsilon^2} \log(\frac{2}{\delta}) \rceil$ iid samples from $\text{Unif}(S \times \mathcal{P}(R))$, **DistEstimate** outputs an ϵ -additive estimate \hat{d}_{S_U} of $d(T, cl_k(S))$ with probability at least $1 - \delta$.

□

Algorithm 3 $\text{UnreachSAT}(\Upsilon)$

- 1: Initialize $\tau \leftarrow 0$.
 - 2: $\tau \leftarrow \text{CryptoMiniSat}(\Upsilon)$.
 - 3: Output $\neg\tau$.
-

Algorithm 4 $\text{Validifier}(P(V, R), S, T, F, \epsilon, \delta, k)$

- 1: Initialize $l \leftarrow \lceil \frac{1}{2\epsilon^2} \log(\frac{2}{\delta}) \rceil$, $\Upsilon \leftarrow \{\}$, $i, j \leftarrow 0$, $D_T \leftarrow \emptyset$, $\tau, d_v \leftarrow 0$.
 - 2: $D_T \leftarrow l$ iid samples from T .
 - 3: **for** $\sigma_i \in D_T$ **do**
 - 4: **for** $j \in [k]$ **do**
 - 5: Construct the formula $\Upsilon = S \wedge F^j \wedge \sigma_i$.
 - 6: $\tau \leftarrow \tau \vee \text{UnreachSAT}(\Upsilon)$.
 - 7: **end for**
 - 8: $d_v \leftarrow (d_v + \frac{\tau}{l})$.
 - 9: $\tau \leftarrow 0$.
 - 10: **end for**
 - 11: Output d_v .
-

1.1.1 Analysis of Validifier

Motivation: The candidate T generated by the decision tree learner **TreeLearner** might contain some program states which are actually not reachable from S in k iterations of the program P , i.e., those states which do not belong to the true k -step invariant $cl_k(S)$. These program states which are not in the true k -step invariant but have been learnt by **TreeLearner** need to be penalised. Hence, we introduce a new weight function in order to quantify how good or bad our candidate T overapproximates the true k -step invariant $cl_k(S)$. This weight function can be formally defined as follows :-

Definition 1.1.3. Over-approximating weight of a given candidate T : Given a probabilistic program P and a set S of initial states, the k -step invariant or the k -iterative closure of P with respect to S is given by the set $cl_k(S) = \{\sigma \in \{0,1\}^{|V|} \mid \sigma \text{ is reachable from } S \text{ in at most } k \text{ iterations of } P\}$. Given a candidate T , the overapproximating weight function for T can be defined as :

$$w(T) = \Pr_{\sigma \sim T}[\sigma \notin cl_k(S)]$$

Claim: Given a probabilistic program P , a set S of initial states, the number of iterations k , a candidate T expressed in CNF, parameters $\epsilon, \delta \in (0, 1]$, **Validifier** outputs an estimate $\hat{w}(T)$ of the weight function $w(T)$ with the following guarantees:

$$\Pr[|\hat{w}(T) - w(T)| \leq \epsilon] \geq (1 - \delta)$$

Moreover, the objective of **Validifier** is to generate a set of negative counterexamples, i.e., those states which have been learnt by the candidate but they do not actually belong to the true k -step invariant $cl_k(S)$. The over-approximating weight function precisely penalises these kind of states for a given candidate T .

Proof. **Description of UnreachSAT:** **UnreachSAT** takes in a CNF formula Υ and returns 1 if Υ is unsatisfiable and 0 otherwise.

Description of Validifier:

Line 2 samples l states from the candidate T almost uniformly at random (using an almost-uniform sampler **CMSGen**) from the witnesses of T to obtain a sample set $D_T = \{y_1, y_2, \dots, y_l\}$, where for each $i \in [l]$, y_i is the output state whose reachability needs to be checked for any $j \in [k]$ iterations of the program P starting from S .

Line 5 constructs for every $j \in [k]$, the j -step reachability formula from S for the sampled output state y_i , denoted by $\Upsilon = S \wedge F^j \wedge y_i$, where F^j is the CNF formula which is satisfiable by all the valid j -length runs of the program S .

Thus, the CNF formula Υ is satisfiable if and only if there exists a valid j -length transition of the program P starting from some state in S and ending up in the final state y_i , i.e., if y_i is reachable in exactly j iterations of P from S . Hence, given a final state y_i , if the collection of CNF formulas $\{S \wedge F^j \wedge y_i\}_{j=1}^k$ is unsatisfiable, then we can conclude that $y_i \notin cl_k(S)$.

Now, let's define the following event for each state $y_i \in D_T$; $i \in [l]$ and $j \in [k]$:

$$E_i : \text{The collection of reachability formulas } \{S \wedge F^j \wedge y_i\}_{j=1}^k \text{ is unsatisfiable.}$$

Thus, the event E_i holds if the sampled final state y_i is reachable in some $j \in [k]$ iterations of P , starting from S . Note that the collection of events $\{E_i\}_{i=1}^l$ are mutually independent. Next, we define indicator random variables for these events as $\mathbf{1}_{E_i}$. Define the statistic $\hat{w}(T) = \frac{1}{l} \sum_{i=1}^l \mathbf{1}_{E_i}$. Line 8 updates the estimate $\hat{w}(T)$ based on the check for non-witness (Line 6) of the collection of formulas $\{S \wedge F^j \wedge y_i\}_{j=1}^k$ generated by the inner loop (Line 5).

We observe that $\mathbb{E}[\hat{w}(T)] = \Pr_{y \sim_U T}[y \text{ is reachable from } S] = w(T)$ and thus, $\hat{w}(T)$ is an unbiased estimator of the quantity $w(T)$. Applying additive Chernoff bound given an error parameter $\epsilon \in (0, 1]$, we get

$$\Pr[|\hat{w}(T) - \mathbb{E}[\hat{w}(T)]| \geq \epsilon] = \Pr[|\hat{w}(T) - w(T)| \geq \epsilon] \leq 2e^{-2l\epsilon^2}$$

We want to make this probability go below a certain threshold, given by δ . Thus,

$$2e^{-2l\epsilon^2} \leq \delta \implies l \geq \frac{1}{2\epsilon^2} \log\left(\frac{2}{\delta}\right).$$

This gives us a sample complexity of $O(\frac{1}{\epsilon^2} \log(\frac{2}{\delta}))$. Thus, we can conclude that if we take atleast $\lceil \frac{1}{2\epsilon^2} \log(\frac{2}{\delta}) \rceil$ iid samples from the CNF formula T , **Validifier** outputs an ϵ -additive estimate $\hat{w}(T)$ of the weight function $w(T)$ with probability at least $1 - \delta$. Also, the number of SAT calls required by **Validifier** in order to verify reachability of the sampled final states = $\lceil \frac{k}{2\epsilon^2} \log(\frac{2}{\delta}) \rceil$.

□

Problem definition:

Problem 1.1.5. (Approximate learning of the k -step invariant $cl_k(S)$): Given a program probabilistic P defined on program variables V , a set S of initial states and parameters $k \in \mathbb{N}$ for the number of program iterations, $\epsilon, \delta \in (0, 1]$, output a candidate \hat{S}_k for the k -step invariant $cl_k(S)$ such that $d(\hat{S}_k, cl_k(S)) \leq \epsilon$ with probability at least $1 - \delta$.

High-level overview of the algorithm:

Ideally, the objective is to learn the k -step invariant $cl_k(S)$. However, it is extremely hard. In this context, can we atleast approximate $cl_k(S)$? That is, we want to output some \hat{S}_k such that $d(\hat{S}_k, cl_k(S))$ is as small as possible. An informal sketch of the algorithm to approximately learn $cl_k(S)$ via \hat{S}_k is given below. The algorithm runs in k phases and tries to learn $cl_k(S)$ in a BFS manner, i.e., it starts off with S and learns $cl_1(S), cl_2(S), \dots, cl_k(S)$ via the sequence $\hat{S}_1, \hat{S}_2, \dots, \hat{S}_k$.

1. **Phase 1:-** Objective is to learn $cl_1(S)$ starting from S .

- Sample *enough* states from S and run the program P for one iteration to obtain a set of output states.
- Build a labeled dataset D using these output states such that the states which are in S are labeled 0 and 1 otherwise.
- Learn a binary decision tree with *bounded size* (according to the size of the formula we want) on D to output a formula φ .
- If $d(S \vee \varphi, cl_1(S)) \geq \epsilon$, perform *weighted secondary sampling* (counterexample-guided sampling). Extend the dataset D by labeling these newly sampled instances.
- Output $\hat{S}_1 = S \vee \varphi$.
- Theoretical guarantees on how close \hat{S}_1 is to the actual one-step invariant $cl_1(S)$ depends on the number of samples taken and the restriction on the size of φ , which is actually dictated by the size of the decision tree learnt.

2. **Phase i ; $i \in \{2, 3, \dots, k\}$:-** Objective is to learn \bar{S}_i starting from \hat{S}_{i-1} .

- Sample *enough* states from S and run the program P for i iterations to obtain a set of output states.
- Build a labeled dataset D using these output states such that the states which are in \hat{S}_{i-1} are labeled 0 and 1 otherwise.
- Learn a binary decision tree with *bounded size* (according to the size of the formula we want) on D . This decision tree would then correspond to the formula φ such that $d(\hat{S}_{i-1} \vee \varphi, \bar{S}_i)$ is minimized.
- Output $\hat{S}_i = \hat{S}_{i-1} \vee \varphi$.
- Once again, theoretical guarantees on how close \hat{S}_i is to the actual i -step invariant \bar{S}_i depends on the number of samples taken and the restriction on the size of φ , which is actually dictated by the size of the decision tree learnt.

Algorithm 5 $\text{ApproxInv}(P(V, R), S, \epsilon, \eta, \delta, k)$

```
1: Initialize  $t \leftarrow \lceil \text{something} \rceil, D_t \leftarrow \emptyset, D \leftarrow \emptyset, \hat{d} \leftarrow \infty, T \leftarrow \{\}, W \leftarrow S$ .
2: for  $j \in [k]$  do
3:    $D_t \leftarrow t$  iid samples from  $\text{Unif}(S \times \mathcal{P}(R))$ .
4:    $D \leftarrow \text{BuildDataset}(P(V, R), D_t, W, j)$ .
5:    $T \leftarrow \text{TreeLearner}(D)$ .
6:    $T \leftarrow \text{Validifier}(P(V, R), W, T)$ .
7:    $\hat{d} \leftarrow \text{DistEstimate}(P(V, R), S, T, \frac{\epsilon}{k}, \delta, j)$ 
8:   while  $\hat{d} \leq \frac{\eta}{k}$  do
9:      $D \leftarrow D \cup \text{SecondarySampler}(P(V, R), S, T)$ 
10:     $T \leftarrow \text{TreeLearner}(D)$ .
11:     $T \leftarrow \text{Validifier}(P(V, R), W, T)$ .
12:     $\hat{d} \leftarrow \text{DistEstimate}(P(V, R), S, T, \epsilon, \delta, j)$ 
13:   end while
14:    $W \leftarrow T$ .
15:    $D_t \leftarrow \emptyset$ .
16:    $D \leftarrow \emptyset$ .
17: end for
18: Output  $W$ .
```

Algorithm 6 $\text{BuildDataset}(P(V, R), D_t, W, j)$

```
1: Initialize  $t \leftarrow |D_t|, D \leftarrow \emptyset, \tau \leftarrow 0$ .
2: for  $i \in [t]$  do
3:   With  $(x_i, R_i) \in D_t$  as initial state, run the program  $P$  for  $j$  iterations to obtain an output state  $y_i$ .
4:    $\tau \leftarrow \text{IsNotWitness}(W, y_i)$ 
5:    $D \leftarrow D \cup (y_i, \tau)$ 
6: end for
7: Output  $D$ .
```

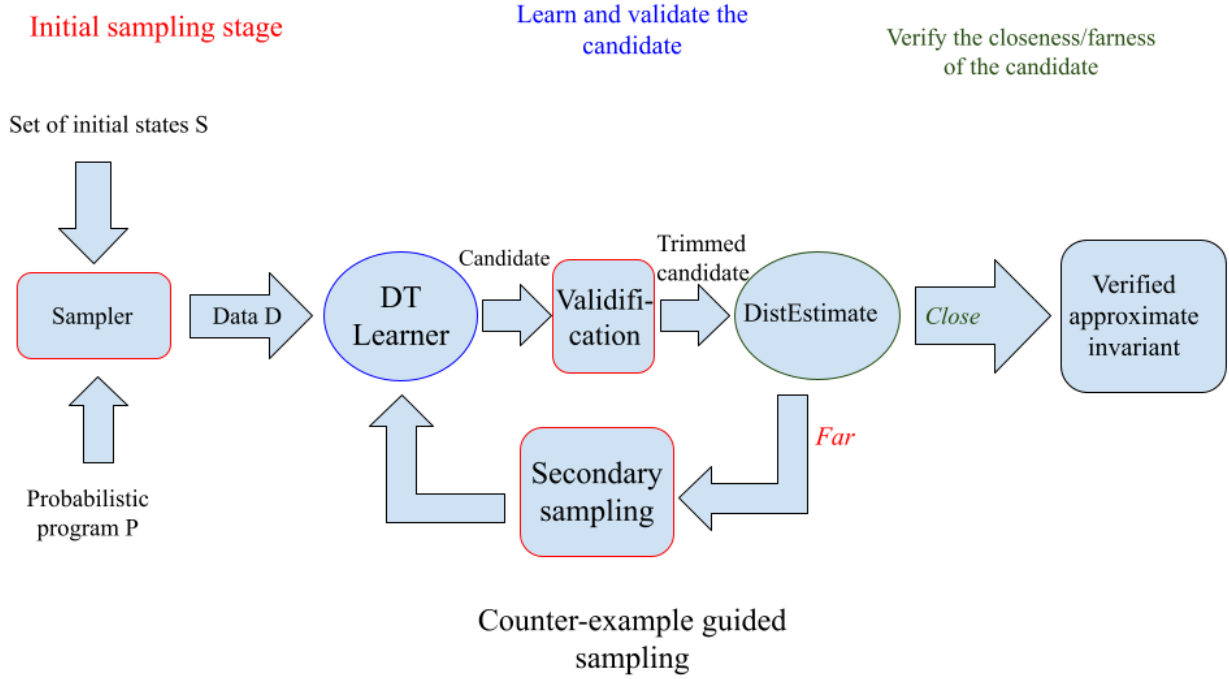
Algorithm 7 $\text{TreeLearner}(D)$

```
1: Initialization.
2: Split-n-Build
3: Prune
4: Output-
```

Algorithm 8 $\text{SecondarySampler}(P(V, R), S, T)$

```
1: Initialization.
2:
3: Output-
```

Figure 1.1: Sketch of ApproxInv



1.1.2 Notes on ApproxInv:

- **BuildDataset:** This subroutine is used for building a labeled dataset for the decision tree learner `TreeLearner`.
- **IsNotWitness:** This subroutine takes in a CNF formula T and an assignment w of its variables and returns 1 if w is not a witness for T and vice-versa.
- **TreeLearner:** This subroutine is used to learn a candidate invariant based on the labeled dataset returned by `BuildDataset`.
- **Validifier:** This subroutine takes in a candidate formula T and a set W of states and trims T by deleting all those states in T which are not reachable in one step from W .
- **ProgCNF:** It takes in a program P and generates a CNF formula F^P corresponding to the valid runs of the program P .
- **DistEstimate:** Description given in Theorem 1.1.4.
- **SecondarySampler:** This subroutine is meant for counterexample-guided sampling in the CEGIS loop.

#samples	100	500	1000	10000
#times test performed	50	50	50	50
d_{\min}	0.11	0.134	0.163	0.1829
d_{\max}	0.25	0.20	0.206	0.1962
d_{mean}	0.1731	0.17296	0.18312	0.1896
d_{std}	0.0284	0.01472	0.01042	0.00313

Table 1.1: Performance of `DistEstimate` on a toy program (Ex9: 5 program variables, 4 internal random variables)

1.1.3 Problems in ApproxInv:

- `TreeLearner` might learn a small-sized formula T such that there exists some state $s \notin cl_k(S)$ but s is a witness for T . In that case $d(T, cl_k(S)) \rightarrow \infty$. *This depends on the pruning scheme of the full decision tree.*

Arnab: `Validifier` can probably be used to trim these spurious states.

- We can do the reachability check while sampling.

1.2 Experimental evaluations on `DistEstimate`: